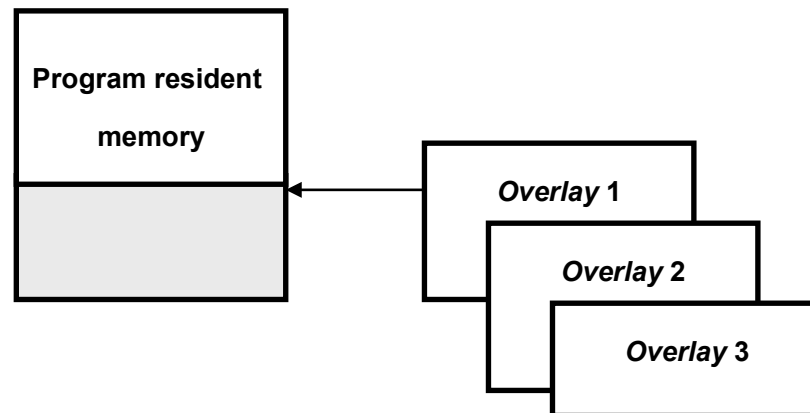


Memory management


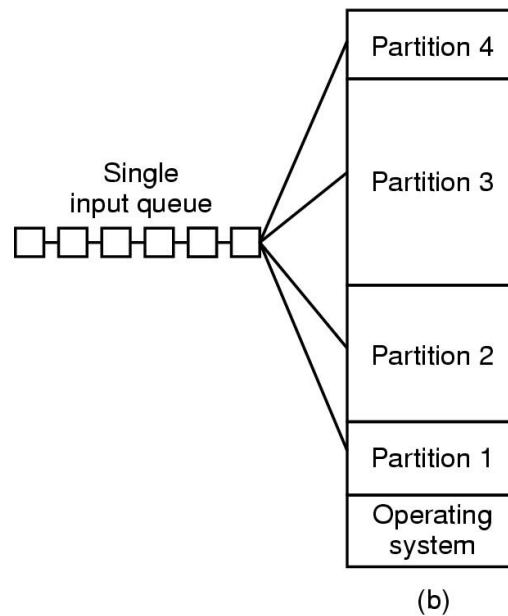
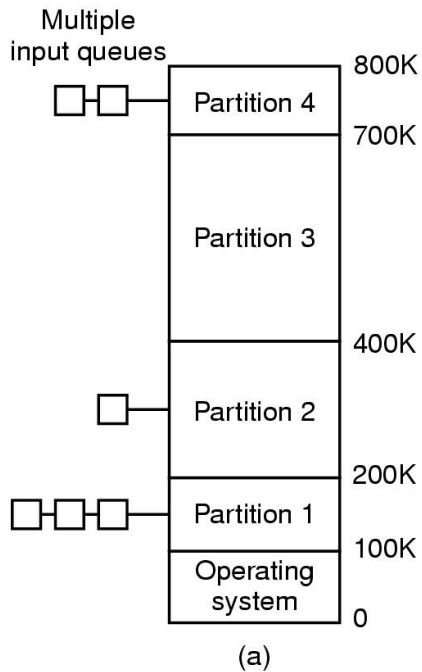
- Single process
 - All memory assigned to the process
 - Addresses defined at compile time
- Multiple processes. How to:
 - assign memory
 - manage addresses?
 - manage relocation?
 - manage program grow?

manage program grow?

- overlay
 - Statically defined memory zone
 - loaded on demand
 - By the program
 - may be replaced



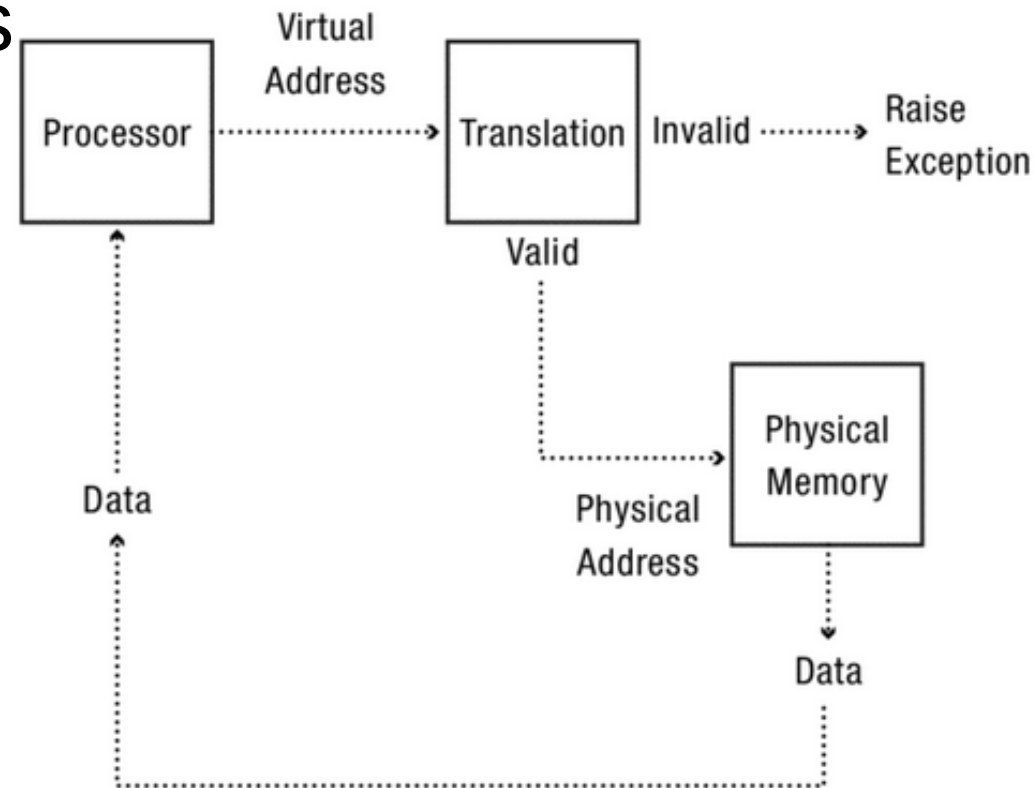
Memory management



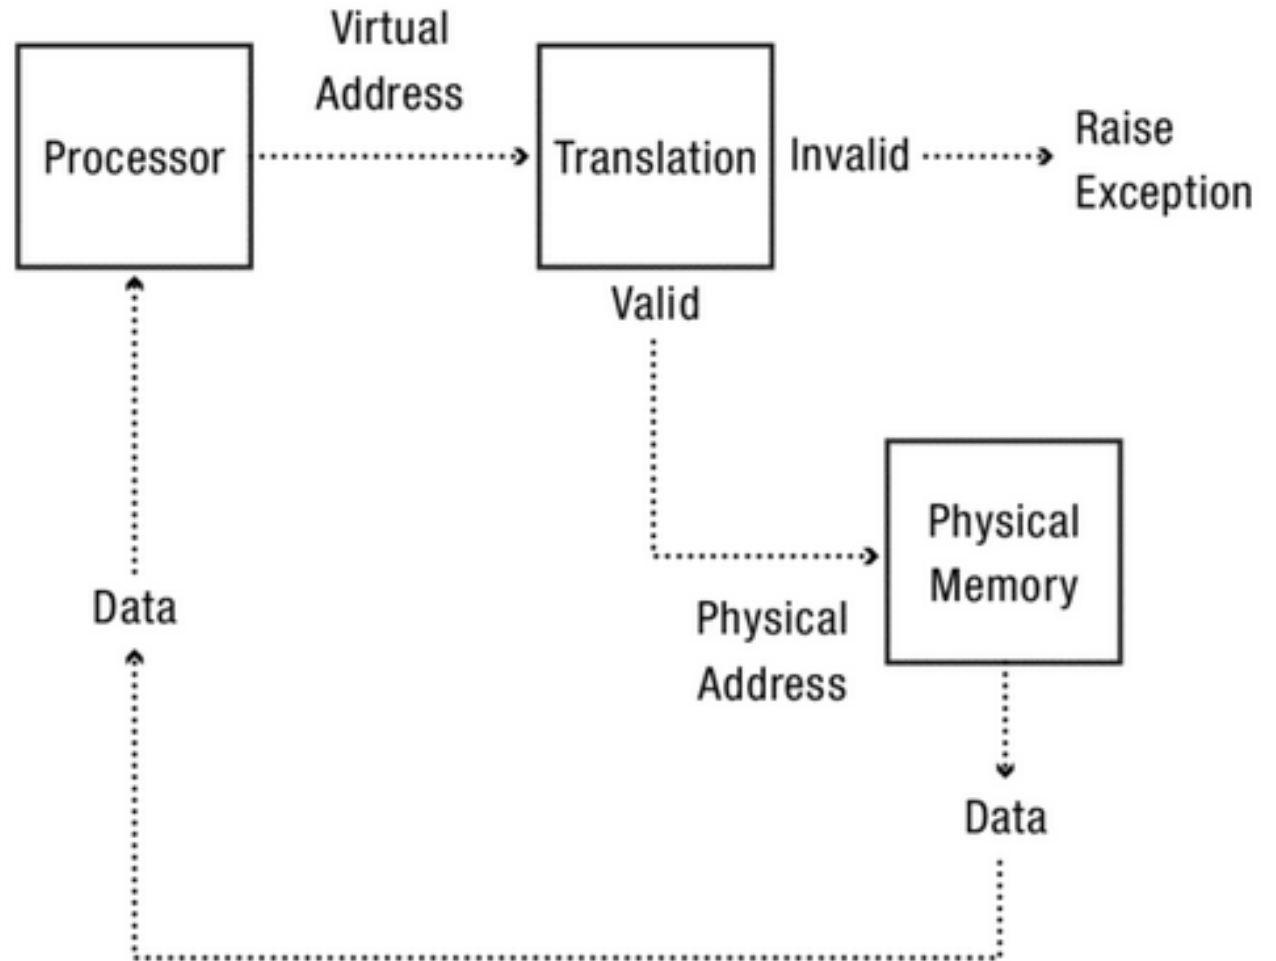
for(i = 0; i < 1000; i++){	.L3:
vect[i] = c;	movl -4(%rbp), %eax
}	cltq
	movzbl -5(%rbp), %edx
	movb %dl, -1008(%rbp,%rax)
	addl \$1, -4(%rbp)
	.L2:
	cmpl \$999, -4(%rbp)
	jle .L3

Virtual memory

- Program addresses
 - are independent on the physical location
- memory manager
 - translates addresses
 - virtual -> physical
 - verifies permissions
- Address Translation



Address Translation Concept




Adress translation

- Process isolation
- Interprocess communication
- Shared code
- Program debugging
- Efficient I/O
- Memory mapped files
- Virtual memory
- Checkpointing
- Process migration
- Information flow control
- DSM

Shared memory

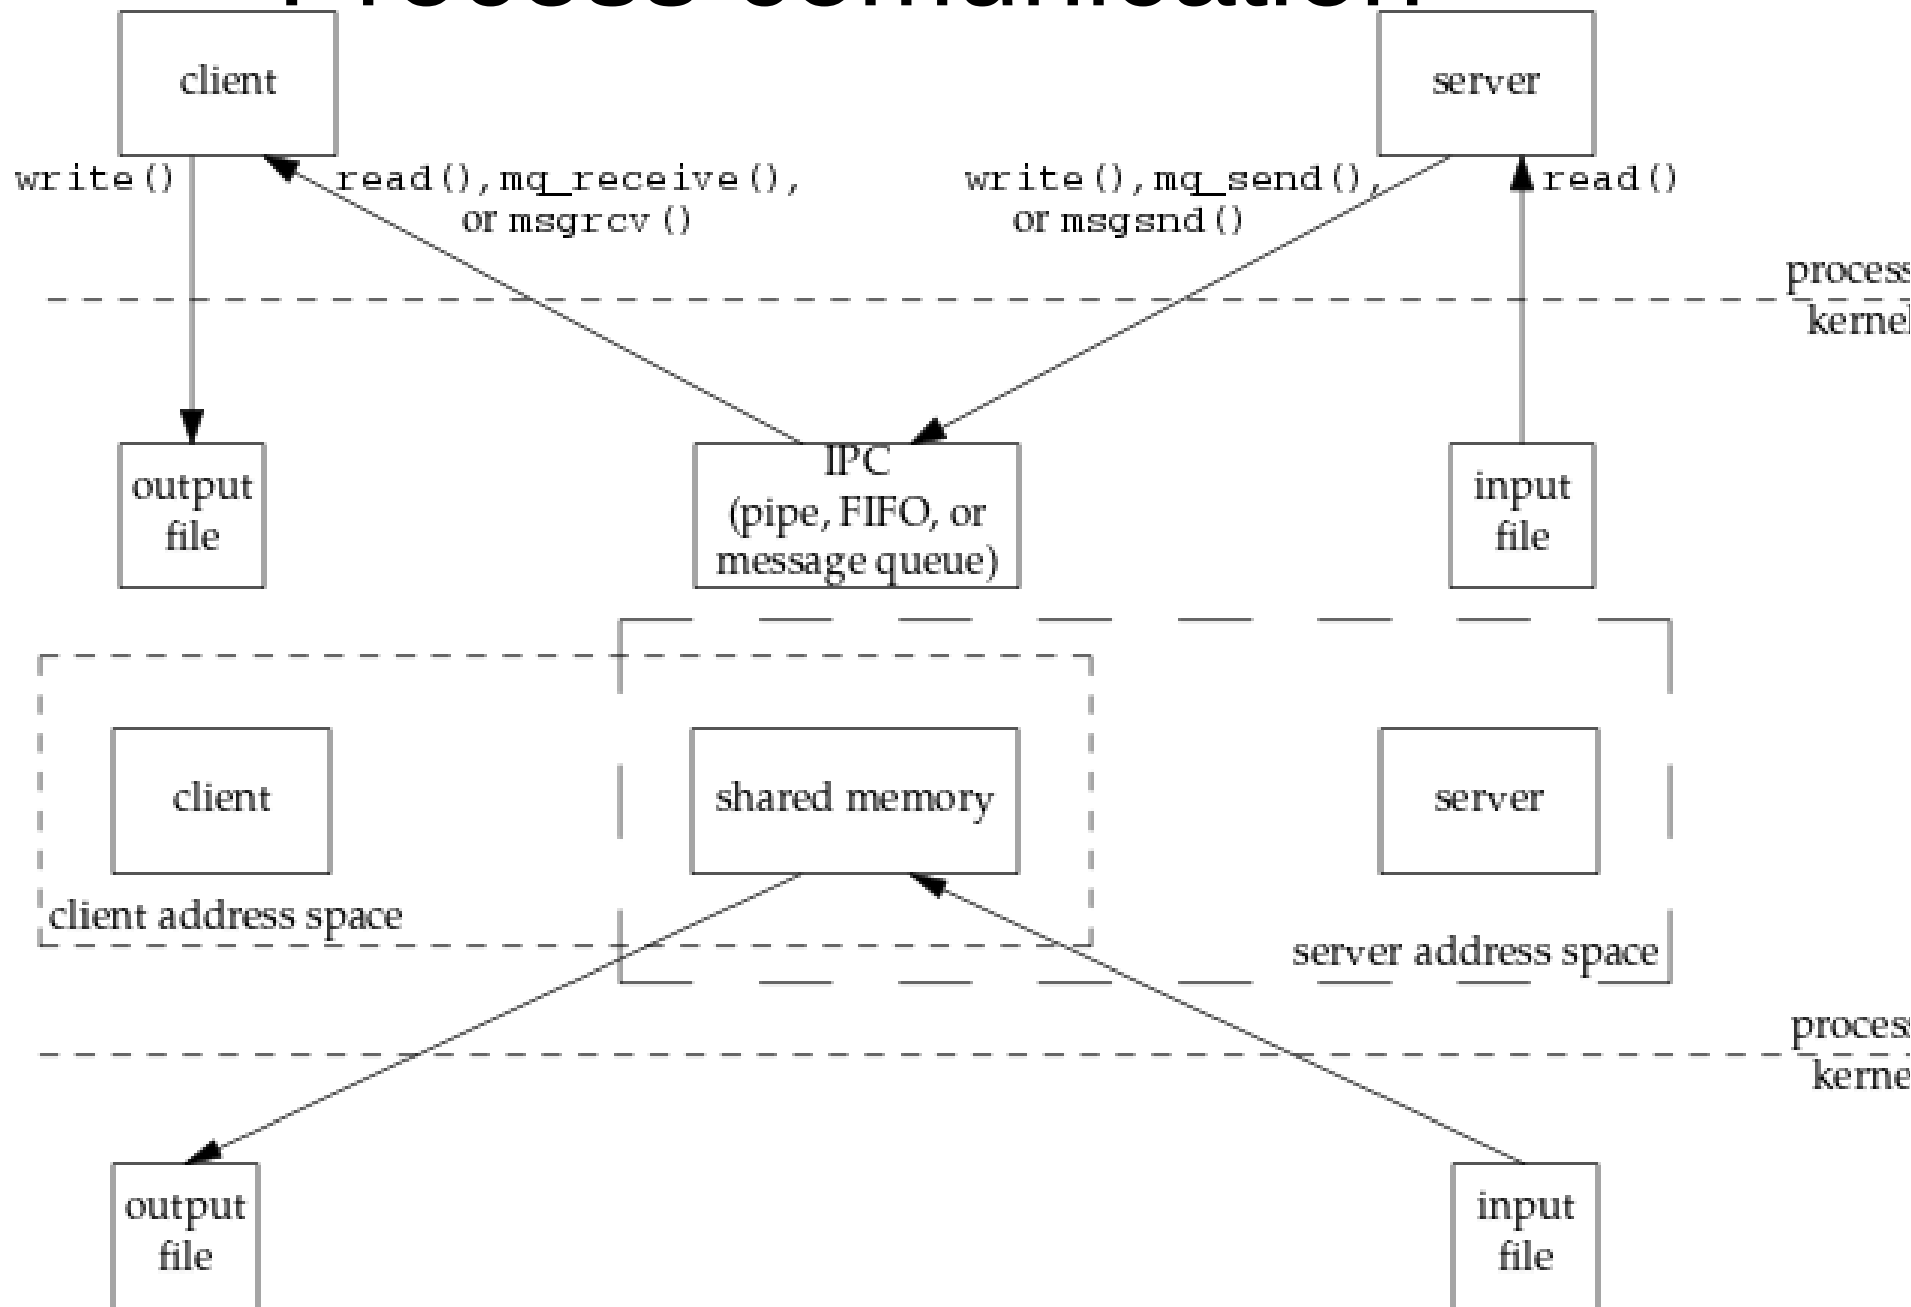
- Processes use regular variables/vector to communicate
 - Available in shared memory
- Should be explicitly created.
 - Each process has its own address space



Shared memory Advantages

- Random Access
 - you can update a small piece in the middle of a data structure, rather than the entire structure
- Efficiency
 - unlike message queues and pipes,
 - copy data between user memory ↔ kernel memory
 - shared memory is directly accessed
 - Shared memory resides in the user process memory
 - Is shared among other processes

Process communication



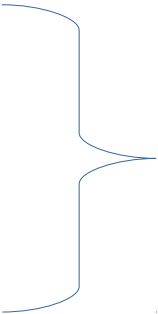
Shared memory Disadvantages

- No automatic synchronization
 - In pipes or message queues
 - Programmer has to provide synchronization
 - Semaphores or signals.
- Pointers are only valid within a given process.
 - Pointer offsets cannot be assumed to be valid across inter-process boundaries.
 - This complicates the sharing of linked lists or binary trees.
- Variables are “produced by the compiler”
 - Names can not be used to access shared memory

Shared memory in *NIX

- System V shared memory
 - Original shared memory mechanism, still widely used
 - Sharing between unrelated processes
- Shared mappings – mmap
 - Shared file mappings
 - Sharing between unrelated processes, backed by filesystem
 - Shared anonymous mappings
 - Sharing between related processes only (related via fork())
- POSIX shared memory
 - Sharing between unrelated processes, without overhead of filesystem I/O
 - Intended to be simpler and better than older APIs

Shared memory in *NIX

- Programming steps
 - Define Shared data structure
 - Creation of memory segment
 - Configuration
 - Assignment to address
 - Access
 - Disconnection
 - Destruction
- 
- In multiple processes

System V shared memory

- Shared memory operations
 - shmget
 - allocates a shared memory segment
 - shmctl
 - allows the user to receive information on a shared memory segment,
 - set the owner, group, and permissions of a shared memory segment,
 - destroy a segment
 - shmat
 - attaches the shared memory segment (identified by shmid) to the address space of the calling process
 - shmdt
 - detaches the shared memory segment (located at the address specified by shmaddr) from the address space of the calling process

System V shared memory

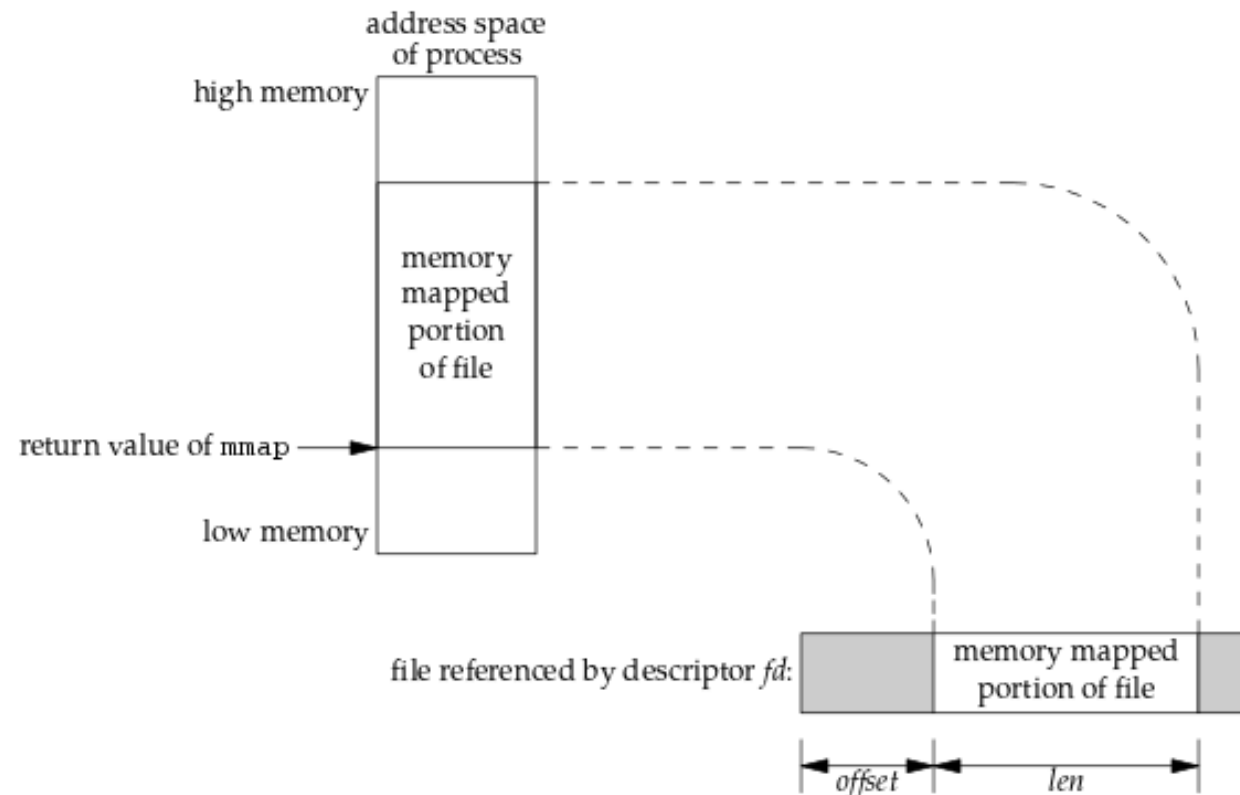
- `int shmget(key_t key, size_t size, int shmflg);`
 - key known by all processes
 - Flags - `IPC_CREAT` | `0666`
- `void *shmat(int shmid, const void *shmaddr, int shmflg);`
 - Shmid – returned by `shmget`
 - shmaddr – `NULL` or other address
 - Shmflg - `SHM_EXEC` `SHM_RDONLY`

System V shared memory

- `char * shm;`
- `key = 5678;`
- `/* Create the segment */`
- `if ((shmid = shmget(key, SHMSZ, IPC_CREAT | 0666)) < 0) {`
- `perror("shmget"); exit(1);`
- `}`
- `/*Now we attach the segment to our data space.*/`
- `if ((shm = shmat(shmid, NULL, 0)) == (char *) -1) {`
- `perror("shmat"); exit(1);`
- `}`
 - Repeated on several processes
 - Key must be known
 - Followed by fork
 - Shm value is shared

Memory mapped files

- Access a file context
 - With memory access operations
 - Assignments e accesses



Memory mapped files

- `void *mmap(void * addr, size_t len ,
int prot, int flags,
int fd, off_t offset);`
 - Prot –
 - PROT_READ PROT_WRITE PROT_EXEC PROT_NONE
 - Flags
 - MAP_SHARED MAP_PRIVATE MAP_FIXED
- `int munmap(void * addr, size_t len);`

Memory access

- After mmap a variable contains a pointer to region
 - Programmer can access that memory as a
 - Pointer to variable
 - Vector
- Is it possible to create linked lists in shared memory?
 - No. mmap in different processes returns different addresses

Memory mapped files

- `Int * ptr;`
- `fd = open(argv[1], O_RDWR | O_CREAT, FILE_MODE);`
- `ptr = Mmap(NULL, sizeof(int), PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0);`
- `Close(fd);`

Mmap and fork

- Open could be avoided
 - Parent and son share the address to memory
 - Memory is shared among processes
- 4.4BSD introduced anonymous memory sharing
 - Flag – MAP_SHARED|MAP_ANON
 - Fd -1
- `ptr = Mmap(NULL, sizeof(int),
PROT_READ | PROT_WRITE,
MAP_SHARED | MAP_ANON,
-1, 0);`

POSIX shared memory

- Mmap
 - File system incurs overhead
- Sharing between unrelated processes, without overhead of filesystem I/O
- Intended to be simpler and better than older APIs
- New function to create memory regions

POSIX shared memory

- Create/opens a shared memory space
 - `fd_mem = shm_open("/myregion", /*region name*/`
 - `O_CREAT | O_RDWR, 0600);`
 - Memory regions are created with size 0
- Assign a size
 - `ftruncate (fd_mem, sizeof(int))`
 - If the object has already been sized by another process, you can get its size with the `fstat` function
- A global region has been created
 - Data stored is kernel persistent
 - But still inaccessible by processes → use `mmap` with `fd_mem`

- `int fd = shm_open(memname,`
- `O_CREAT | O_TRUNC | O_RDWR,`
`0666);`
- `if (fd == -1)`
- `error_and_die("shm_open");`
- `int r = ftruncate(fd, sizeof(int));`
- `if (r != 0)`
- `error_and_die("ftruncate");`
- `int *v_int = mmap(0, sizeof(int),`
- `PROT_READ | PROT_WRITE, MAP_SHARED,`
- `fd, 0);`

Shutdown

- Close the shared memory object
 - `close(fd_mem)`
- Unmap the shared memory object:
 - `munmap (pointer, SHM_SIZE);`
 - The address become free to be used.
- At this stage if other processes open and map a memory region
 - They can access data previously written
- To remove permanently the shared memory object:
 - `shm_unlink (SHARED_MEMORY_NAME);`
 - The object is effectively deleted after the last call to `shm_unlink`

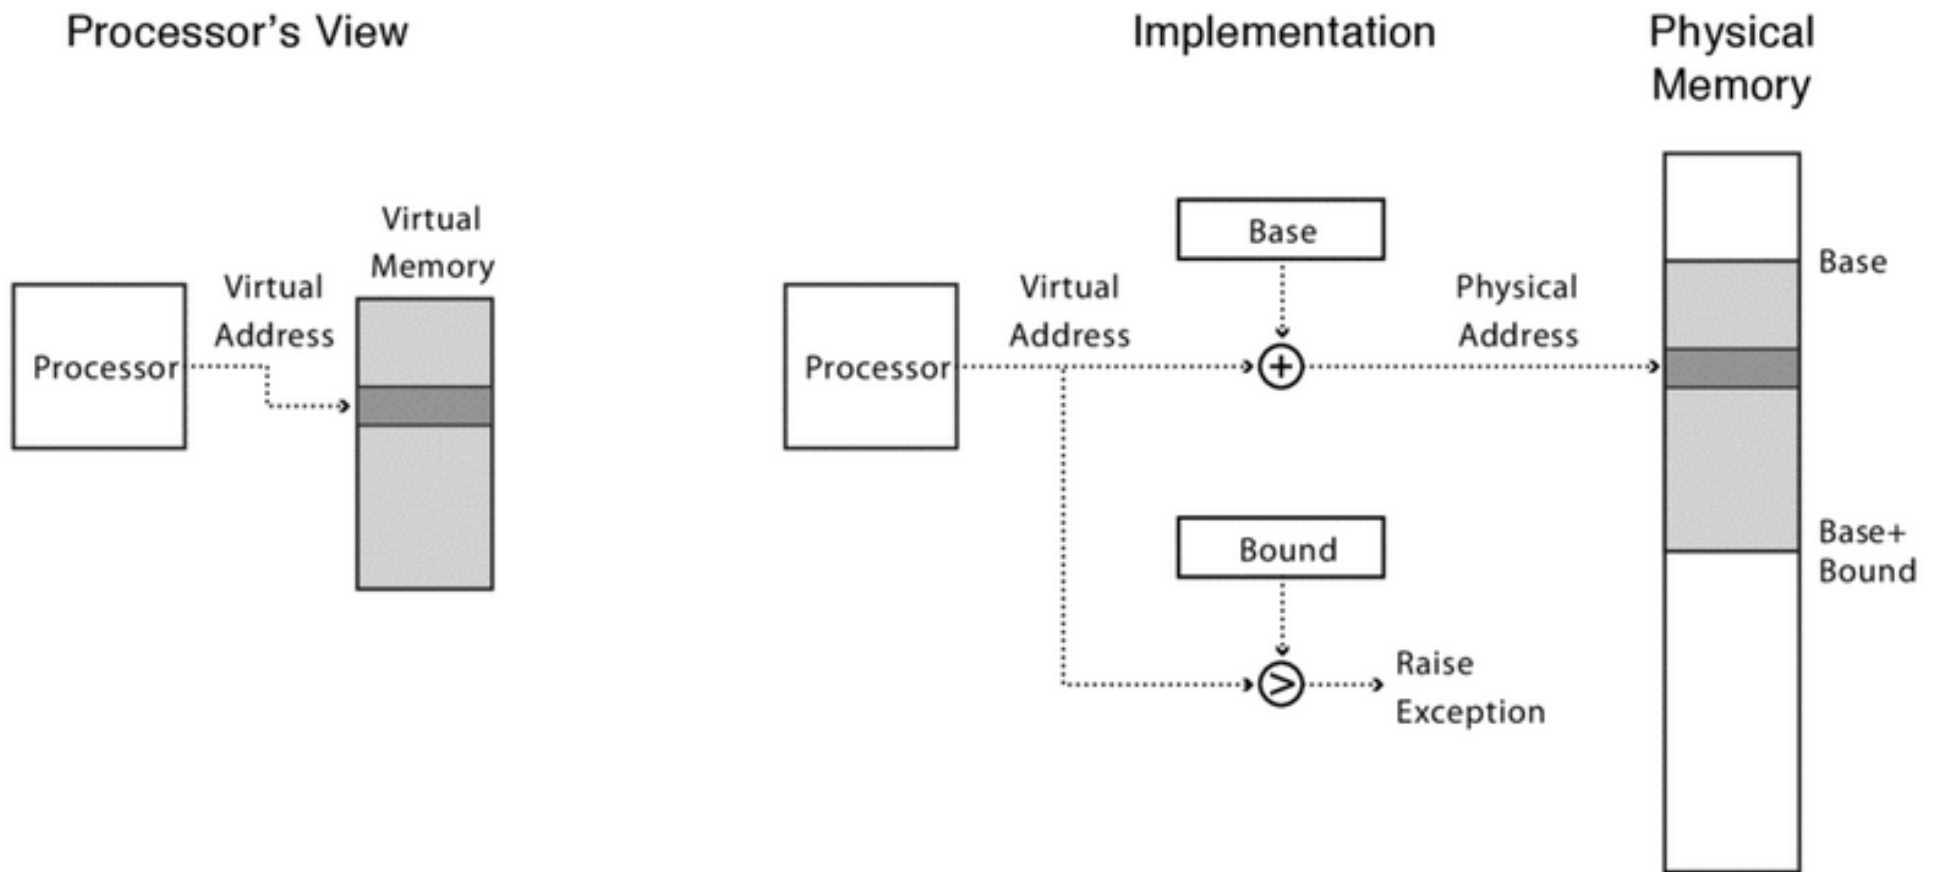
Address Translation

- Flexible Address Translation
 - Base and bound
 - Segmentation
 - Paging
 - Multilevel translation
- Efficient Address Translation
 - Translation Lookaside Buffers
 - Virtually and physically addressed caches

Address Translation Goals

- Memory protection
- Memory sharing
 - Shared libraries, interprocess communication
- Sparse addresses
 - Multiple regions of dynamic allocation (heaps/stacks)
- Efficiency
 - Memory placement
 - Runtime lookup
 - Compact translation tables
- Portability

Virtually Addressed Base and Bounds



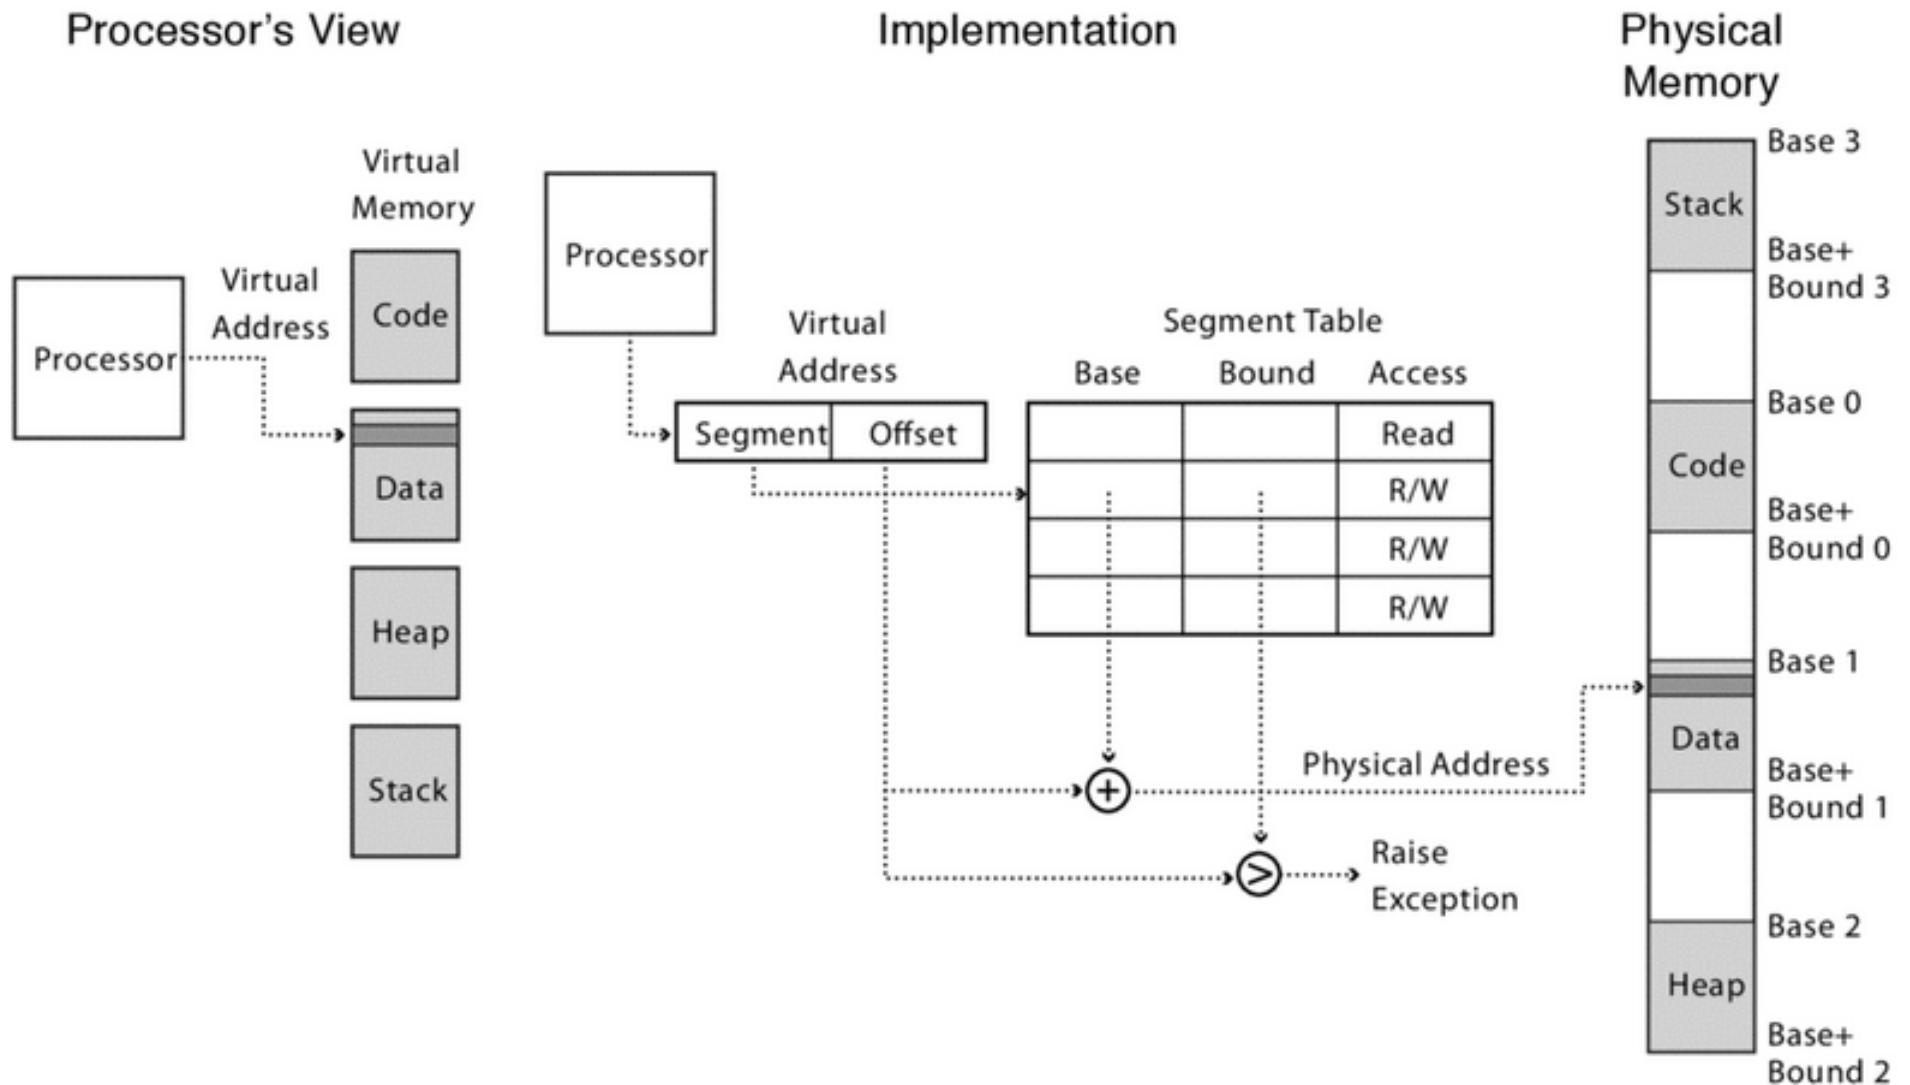
Virtually Addressed Base and Bounds

- Simple and fast
- coarse-grained protection
 - at the process levels
 - impossible to prevent self overriding of code
- Difficult to share memory regions
- Memory needs to be continuous
 - hard to implement dynamic memory
 - stack heap, ...

Segmentation

- Evolution of base and bound
 - one process -> multiple bases+bounds
 - one process -> multiple segments
- Segment: base + bound
 - multiple sizes
 - each segment stored continuously
 - multiple access permissions
 - execute-only (code) or read-write (data)

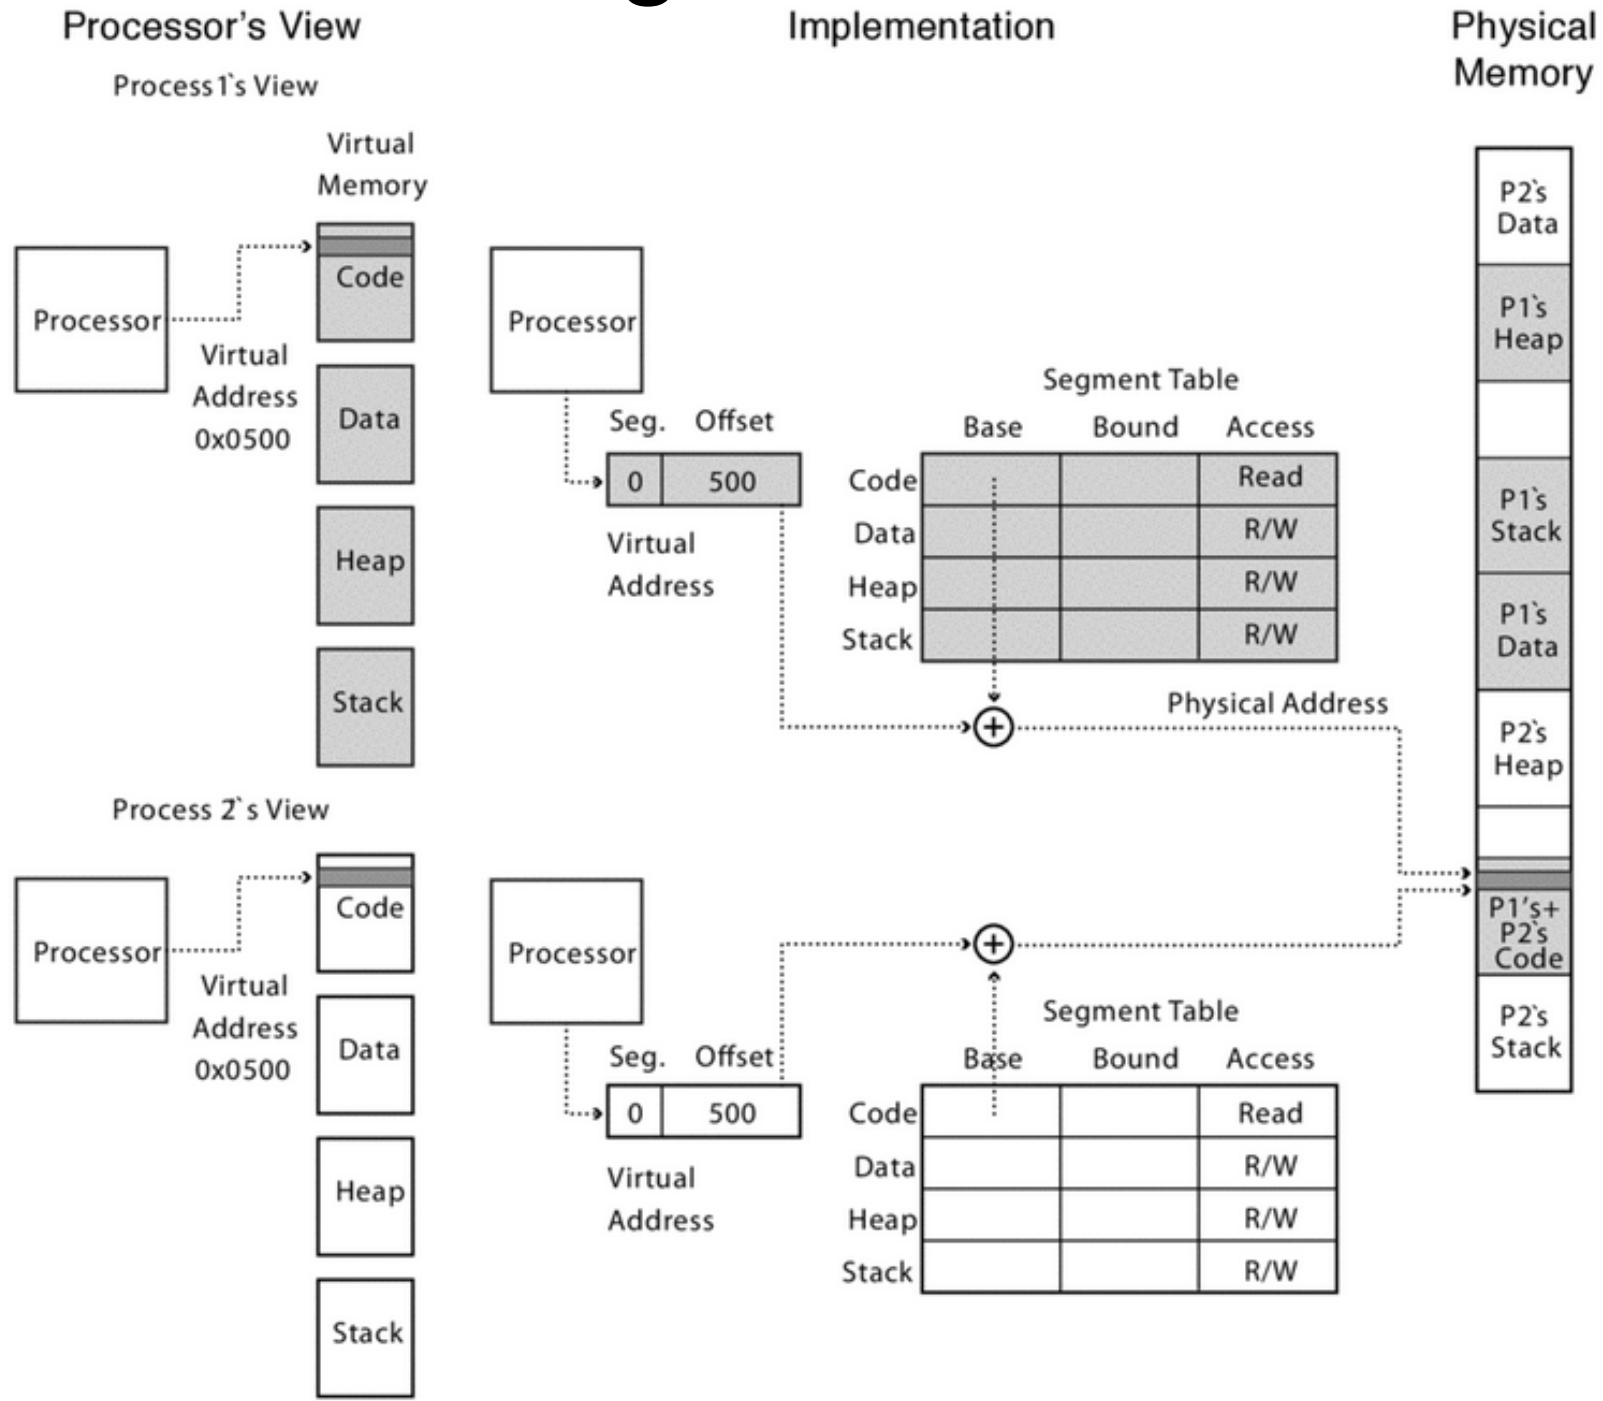
Segmentation



Segmentation

- Segmented memory has gaps
 - program memory is not a single continuous regions
 - each segment start at a different base
- Access outside segments
 - HW generates interrupt
 - Kernel generates segmentation fault
- Addressable space
 - depends on the segment+offset length
-

Segmentation



Segmentation

- Total memory depends
 - base + offset
- Kernel should manage
 - placement of each segment on physical memory
 - permission of access

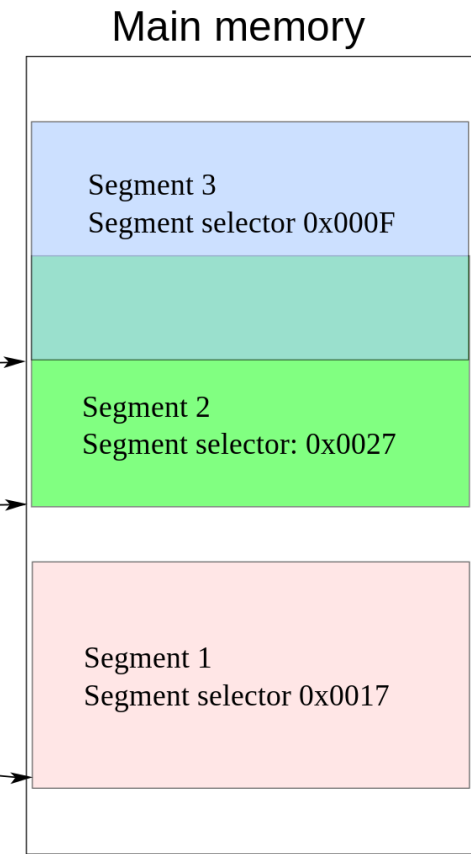
Segmentation

- Overhead of managing
 - Large number of segments
 - Segments of variable size
 - Dynamically growing segments
- Creation of a segment
 - Virtual address (easy)
 - Memory location (difficult)
 - Fragmentation
 - Segments grows

Local Descriptor Table (LDT)

5			
4	0x21430	0xC000	•
3			
2	0x0CEF0	0xA300	•
1	0x28C00	0xFC00	•
0			

Linear base address (BASE) Segment size (LIMIT)



Paged memory

- Alternative to segmentation
 - Fixed sized chunks
 - Page frames
 - Address translation is similar to segmentation
 - Segment table → page table
 - Table entry points to page base
 - No need for bound
- Process still views memory as linear
 - Linear memory (virtual) accesses
 - Multiple accessed pages

Paged memory

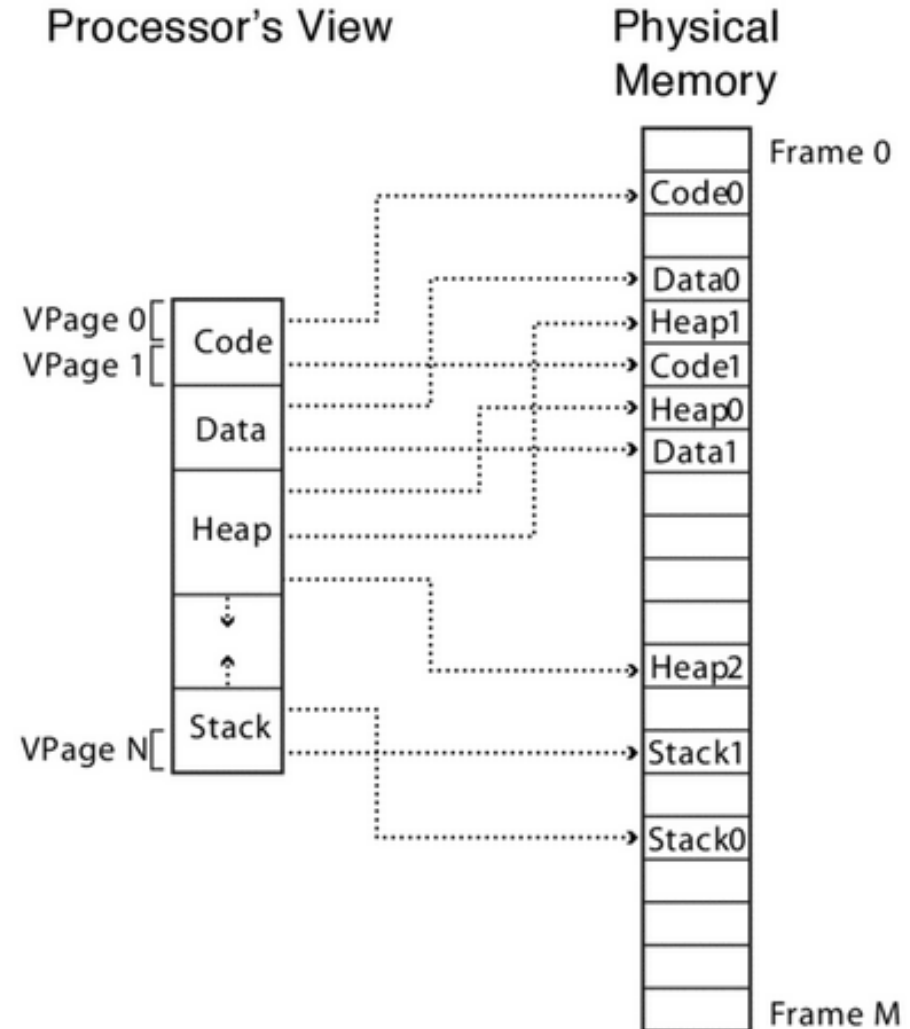
- Big advantage over segmentation
 - Free-space allocation is straightforward
 - Physical memory representation
 - Bitmap (one bit per physical page)
- Memory sharing
 - Setting of pointer in page table

Optimizations

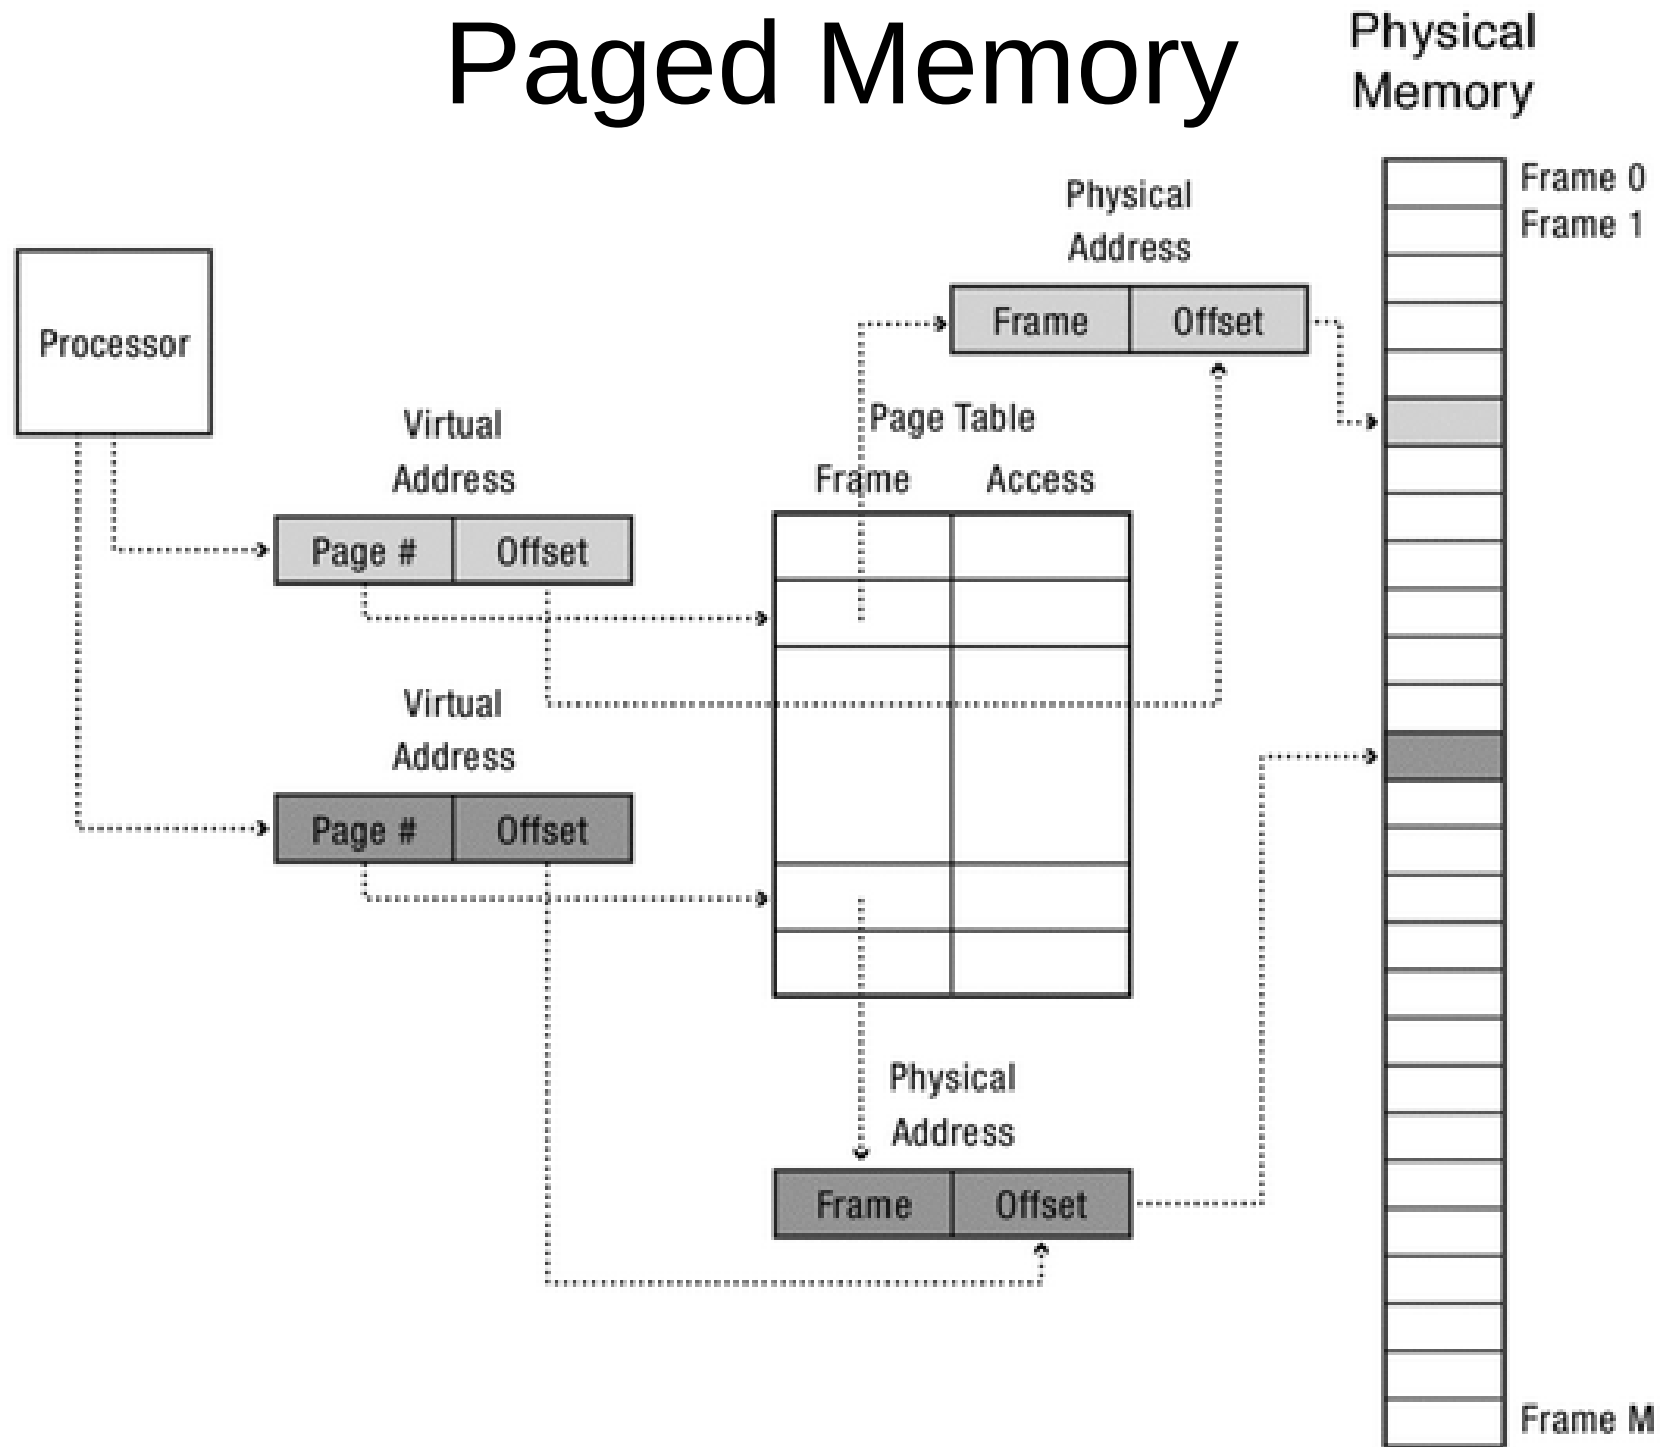
- Segmentation+paged memory
 - Copy on write
 - Zero on reference
 - Data breakpoints
 - Stop program when variable gets changed
- Paged memory
 - Program can execute without all code/data
 - Pages can be loaded while program is executing
 -

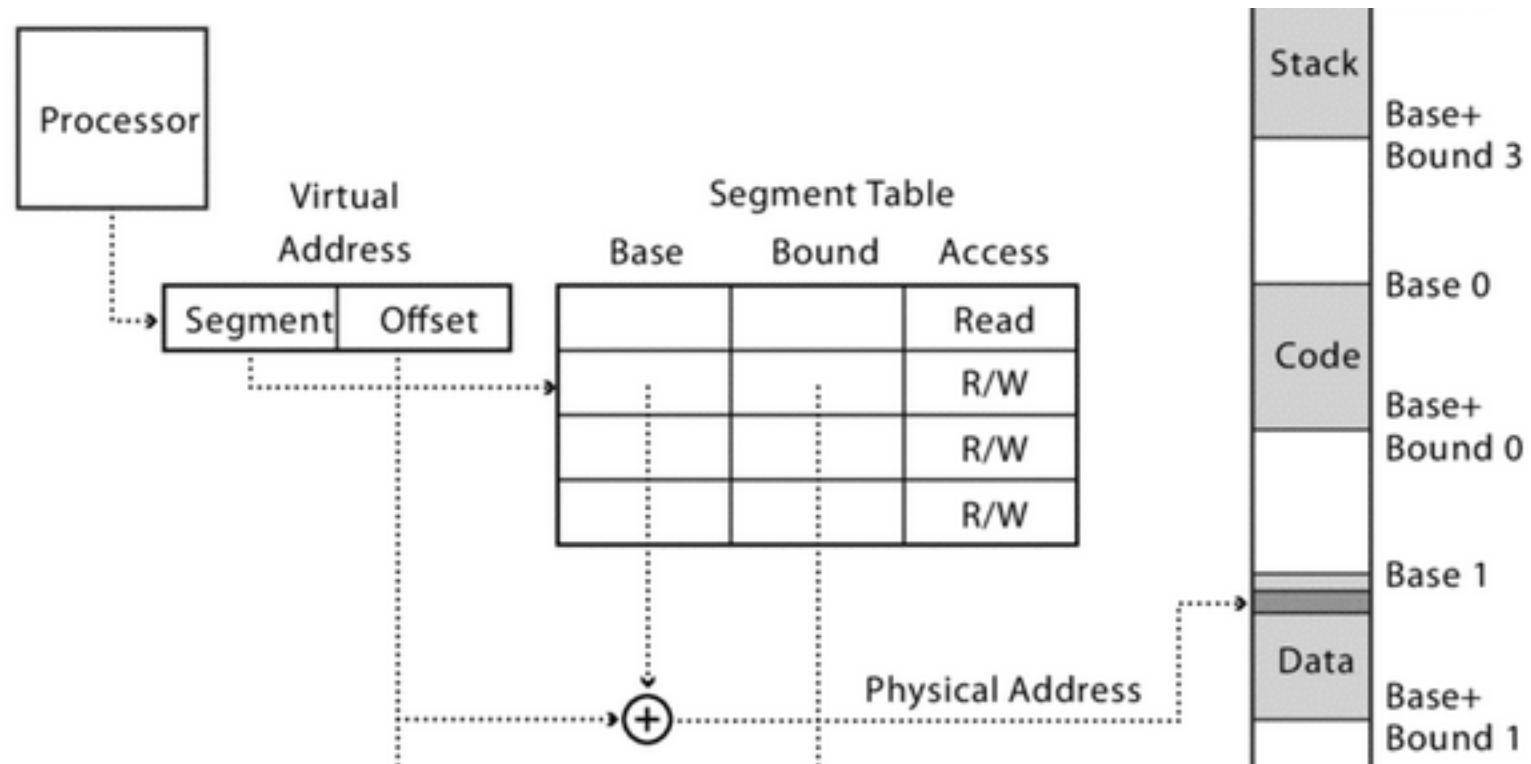
Paged memory

- Downsides
 - Virtual address management is more difficult
 - Stack and heap should be continuous
 - Compiler should guarantee this
 - How to make it work with multiple threads?

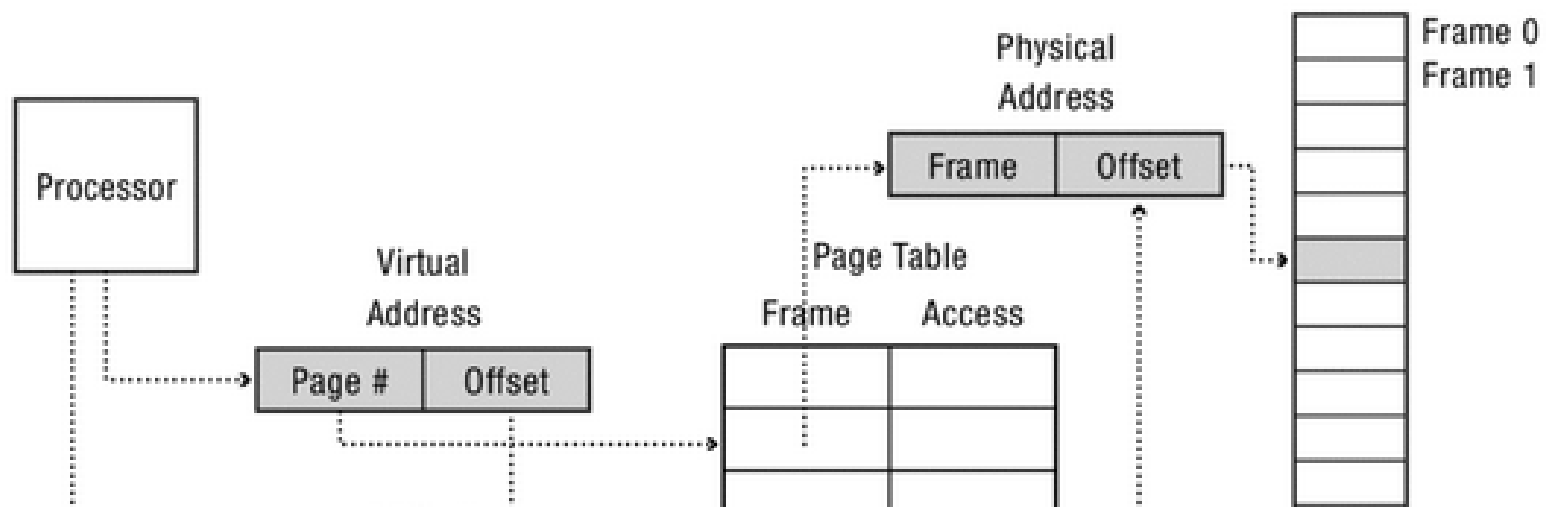


Paged Memory





- Segmentation
- Page



How to define virtual addresses?

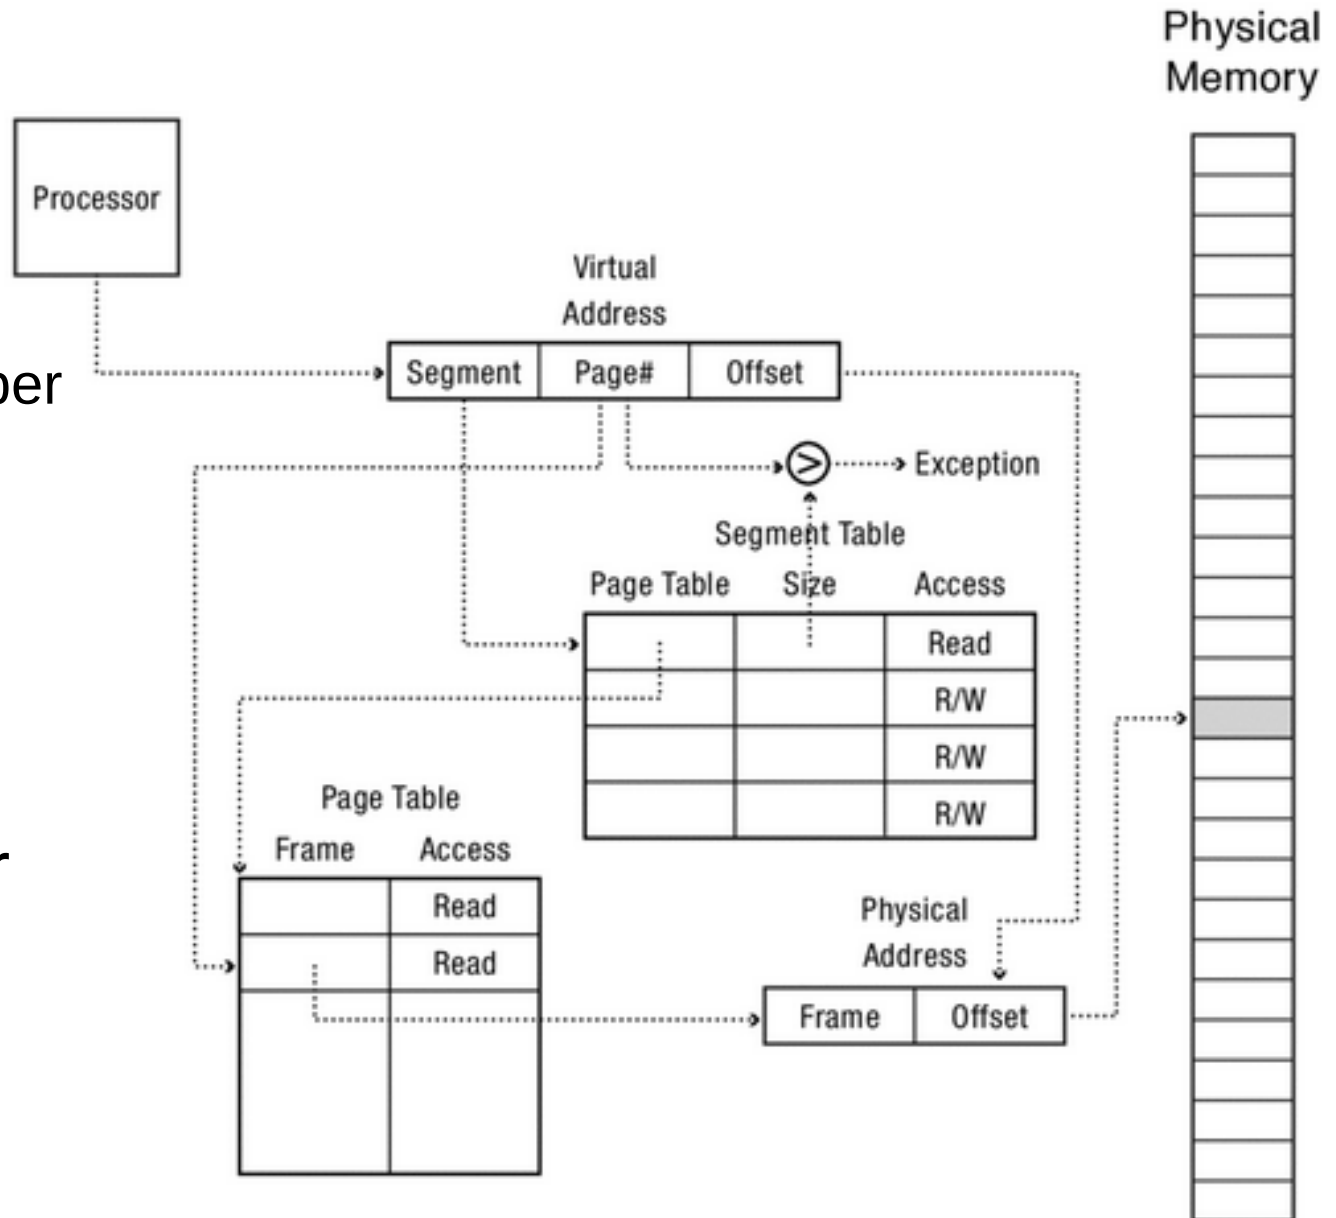
- Segmentation
 - Virtual address \rightarrow segment# + offset
 - Max size of the segment \rightarrow offset
 - Max memory per process \rightarrow segment table size
 - Max real memory \rightarrow base + offset

Where are the tables

- In memory
 - Special page/segment
- Multiple organizations
 - Page segmentation
 - Multi-level paging
 - Multi-level paged segmentation

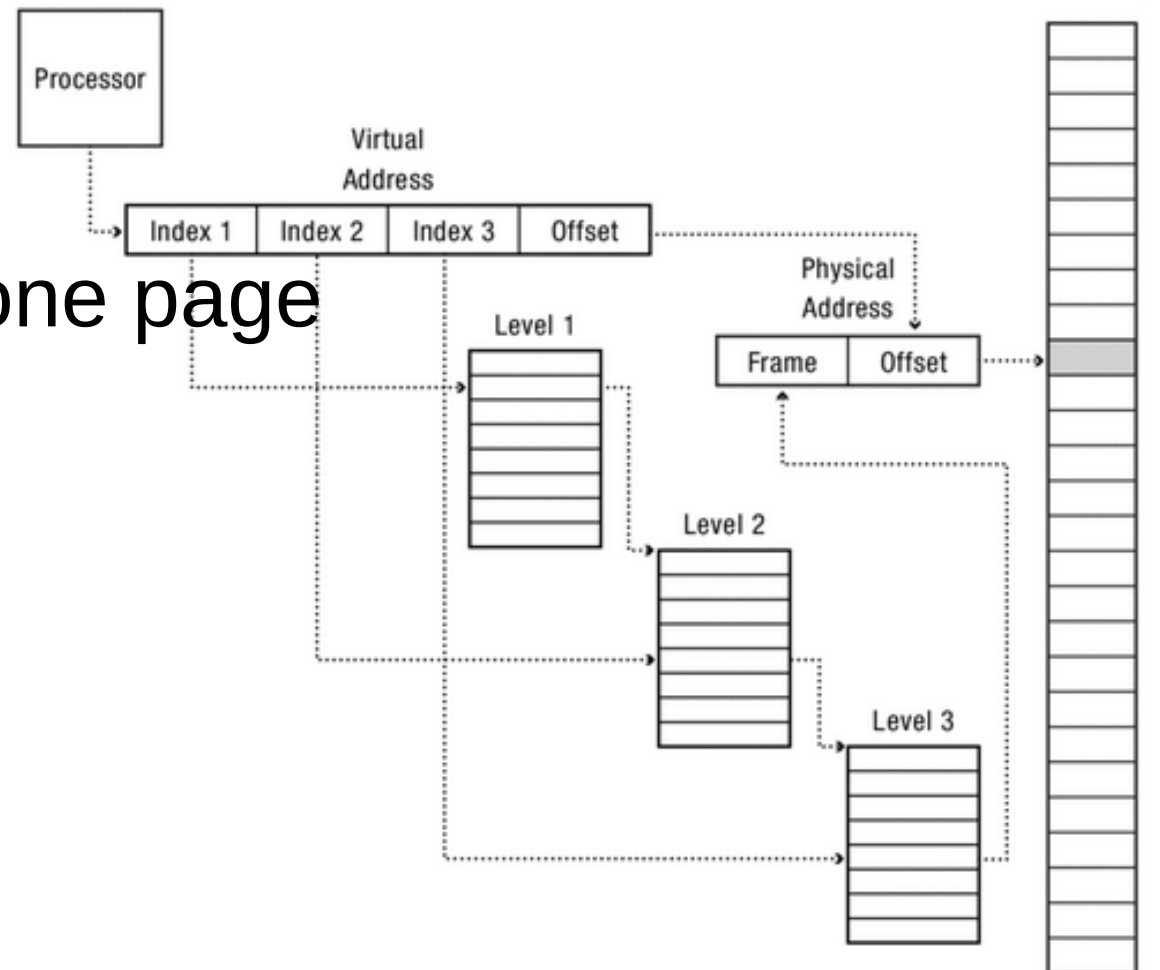
Page segmentation

- 32 bit address
- 4Kb pages
 - 10 bits
 - Segment number
 - 10 bits
 - Page number
 - 12 bits
 - Page offset
- 1 page table per segment
 - 4 bytes
 - 1K entries



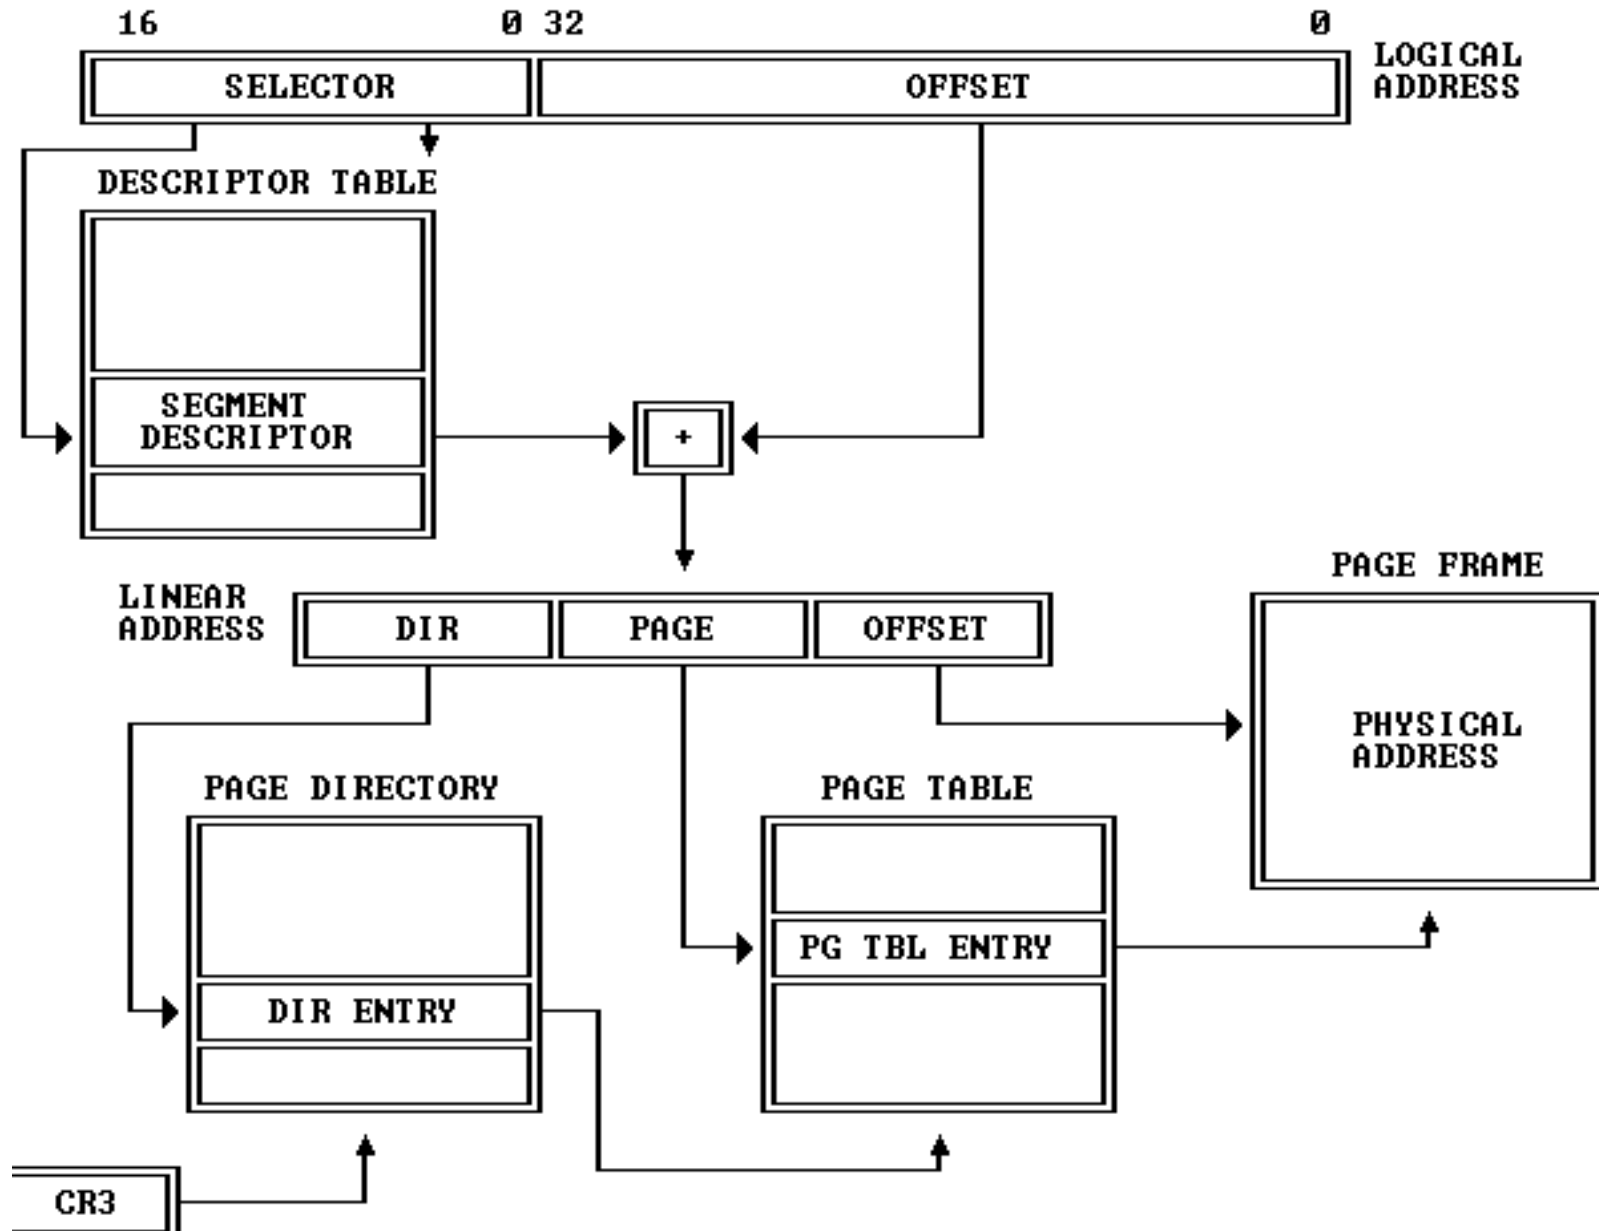
Multi-level paging

- Virtual address has 4 components
 - 3 indexes
 - Offset
- Each tabel fits into one page



Multi-level paged segmentation (X86)

figure 5-12. 80386 Addressing Mechanism



32 bit - X86

- 2 level page table within a segment
 - 10 bits – page directory
 - 10 bit – second level page table
 - 12 bit – offset
- 4kb page frame
- Number of 2nd level page table
 - Depends on length of segment
 - Also has permissions

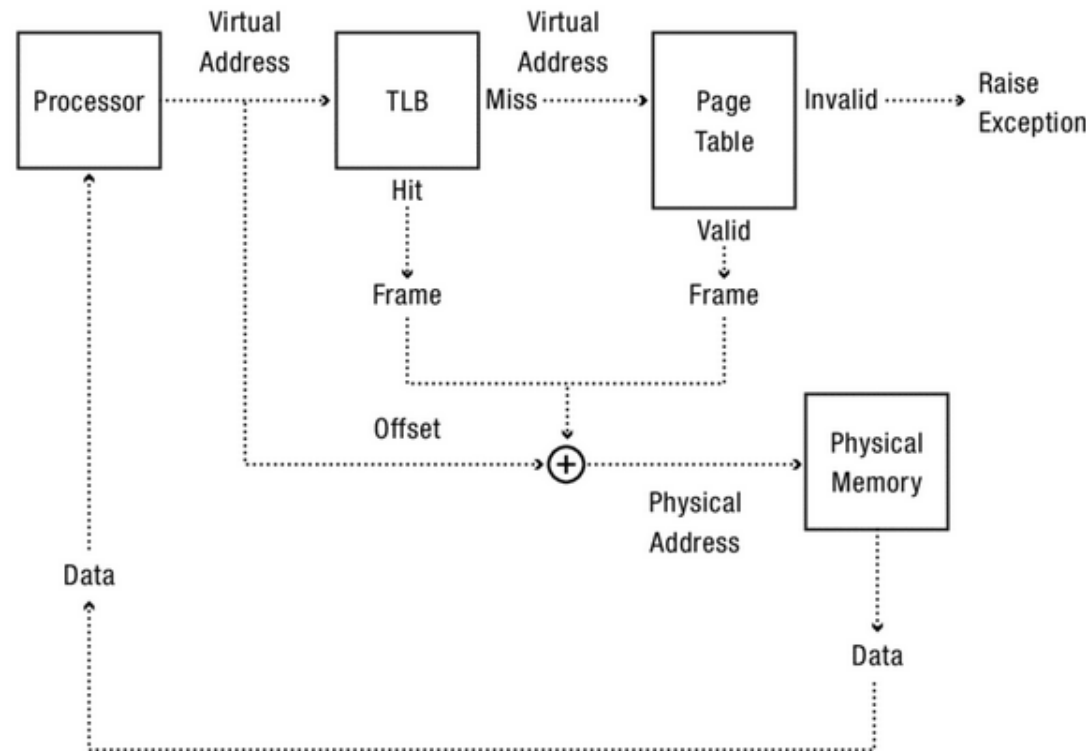
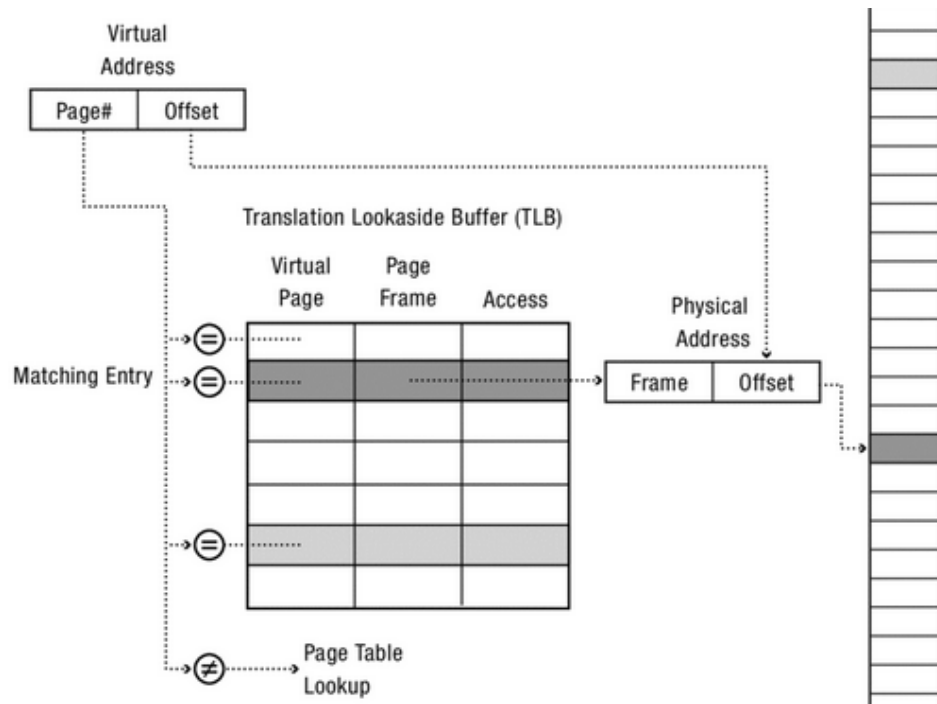
64 bit - X86

- 64 bits maximum address length
 - Optimization addresses only use 48 bit 128 terabytes
- 4 level page table within a segment
- Optimizations
 - Elimination of two levels of page tables
- 4kb page frame
 - 4th level pages → 2Mb
 - Can be overwritten if OS allocates 2Mb contiguous
 - 3rd level pages → 1Gb
 - Can be overwritten if OS allocates 1Gb contiguous
 - Good for server with two terabytes

Efficient Address translation

- Conversion from virtual to real address
 - Requires extra memory accesses
- Memory accesses are sequential
 - Repeated accesses to same page
 - Why not cache last translations?
- Translation look-aside buffers
 - Small HW table caching recent address translation

Translation look-aside buffers



TLB

- Extra HW
- Big savings in memory accesses
- Multiple levels of TLBs
- Requires extra care
 - Multiple processes
 - Different virtual address space
 - Different permissions
 - Multiprocessors
 - Consistency of multiple TLB
 - one per CPU